

# Decorator System: Type-Safe Function Enhancement

## Overview

The decorator system provides **type-safe cross-cutting concerns** using Python's TypeVar and Protocol systems. Unlike traditional decorators that break type checking, these decorators preserve the original function signature for perfect IDE support.

## The TypeVar Advantage

### Traditional Decorator Problem

```
# Traditional decorator loses type information
def my_decorator(func): # No type info!
    def wrapper(*args, **kwargs): # Generic args
        return func(*args, **kwargs)
    return wrapper

@my_decorator
def add(a: int, b: int) -> int: # Type info lost!
    return a + b

# IDE doesn't know 'add' signature anymore
result = add(1, 2) # No type checking, no autocomplete
```

### TypeVar Solution

```
from typing import Callable, TypeVar, Any, cast
from functools import wraps

# Define TypeVar bound to Callable
F = TypeVar("F", bound=Callable[..., Any])

def my_decorator(func: F) -> F:
    """Decorator that preserves function signature."""

    @wraps(func)
    def wrapper(*args: Any, **kwargs: Any) -> Any:
        print(f"Calling {func.__name__}")
        return func(*args, **kwargs)

    # Cast to F preserves original signature
    return cast(F, wrapper)

@my_decorator
def add(a: int, b: int) -> int: # Type info preserved!
    return a + b
```

```
# IDE knows exact signature!
result = add(1, 2) # ✓ Type checks
result = add("a", "b") # X Type error caught
```

## Available Decorators

### 1. `@validate_context`

**Purpose:** Validate AgentContext before execution

**Returns:** AgentResult with errors instead of raising exceptions

```
@validate_context
def execute(context: AgentContext) -> AgentResult:
    # context is guaranteed valid here
    return AgentResult(...)
```

#### Checks:

- Context is not None
- Context is AgentContext instance
- Context has required fields (task, etc.)

#### Benefits:

- Fail-fast on invalid input
- Consistent error handling
- No exception propagation

---

### 2. `@timed`

**Purpose:** Measure execution time

**Adds to metadata:** `execution_time_ms`, `execution_timestamp`

```
@timed
def execute(context: AgentContext) -> AgentResult:
    # Expensive operation
    return AgentResult(...)

# Result automatically has:
# metadata['execution_time_ms'] = 42.5
# metadata['execution_timestamp'] = '2026-01-22T10:30:00'
```

#### Use Cases:

- Performance monitoring

- Identifying slow skills
  - SLA tracking
  - Billing/metering
- 

### 3. @logged

**Purpose:** Structured logging of execution

**Logs:** Entry, exit, status, errors

```
@logged
def execute(context: AgentContext) -> AgentResult:
    return AgentResult(...)

# Logs:
# INFO: → Entering execute [task=analyze_data]
# INFO: ← Exiting execute: SUCCESS [42.5ms]
```

#### Benefits:

- Automatic audit trail
  - Debugging assistance
  - Performance visibility
  - No manual logging needed
- 

### 4. @cached(ttl\_seconds)

**Purpose:** Cache results to avoid redundant computation

**Cache Key:** Generated from `context.task` + `context.parameters`

```
@cached(ttl_seconds=300) # Cache for 5 minutes
def execute(context: AgentContext) -> AgentResult:
    # Expensive computation
    return AgentResult(...)

# First call: executes and caches
# Second call (within 5 min): returns cached result
```

#### Metadata Added:

- `cache_hit`: True/False
- `cached_at`: ISO timestamp

#### Use Cases:

- Expensive computations
- Rate-limited APIs

- Database queries
  - External service calls
- 

## 5. `@retry(max_attempts, delay_seconds)`

**Purpose:** Retry failed executions with exponential backoff

**Returns:** Only on FAILURE status, not exceptions

```
@retry(max_attempts=3, delay_seconds=1.0)
def execute(context: AgentContext) -> AgentResult:
    # Potentially flaky operation
    return AgentResult(...)

# Attempt 1: fails, wait 1s
# Attempt 2: fails, wait 2s
# Attempt 3: fails, return last failure
```

### Metadata Added:

- `retry_attempt`: Current attempt number
- `retry_needed`: Whether retries occurred
- `all_attempts_failed`: True if all failed

### Use Cases:

- Network calls
  - Transient failures
  - Rate limit handling
  - External API integration
- 

## 6. `@require_params(*params)`

**Purpose:** Validate required parameters exist

**Returns:** FAILURE if any parameter missing

```
@require_params('dataset', 'operation')
def execute(context: AgentContext) -> AgentResult:
    # Guaranteed to have these parameters
    dataset = context.parameters['dataset']
    operation = context.parameters['operation']
    return AgentResult(...)
```

### Error Details:

- `missing_params`: List of missing parameters
- `required_params`: List of required parameters

- `received_params`: List of provided parameters

### Use Cases:

- Input validation
  - Fail-fast on missing data
  - Clear error messages
  - Self-documenting requirements
- 

## 7. `@enrich_metadata(**fields)`

**Purpose:** Add custom metadata to results

**Usage:** Tagging, categorization, versioning

```
@enrich_metadata(  
    skill_category='analytics',  
    version='1.0.0',  
    author='data-team'  
)  
def execute(context: AgentContext) -> AgentResult:  
    return AgentResult(...)  
  
# Result automatically has:  
# metadata['skill_category'] = 'analytics'  
# metadata['version'] = '1.0.0'  
# metadata['author'] = 'data-team'
```

### Use Cases:

- Skill categorization
  - Version tracking
  - Team attribution
  - Custom tagging
- 

## 8. `@standard_skill_decorators`

**Purpose:** Apply common decorator stack

**Equivalent to:** `@logged @timed @validate_context`

```
@standard_skill_decorators  
def execute(context: AgentContext) -> AgentResult:  
    # Gets validation, timing, and logging automatically  
    return AgentResult(...)
```

### Benefits:

- Consistent behavior across skills
  - Single decorator for common needs
  - Reduces boilerplate
- 

## Decorator Stacking

Decorators can be stacked for combined functionality:

```
@enrich_metadata(category='analytics') # 5. Add metadata
@cached(ttl_seconds=300)             # 4. Cache results
@retry(max_attempts=3)              # 3. Retry on failure
@require_params('dataset')          # 2. Validate params
@standard_skill_decorators         # 1. Validate, time, log
def execute(context: AgentContext) -> AgentResult:
    return AgentResult(...)
```

**Execution Order** (bottom to top):

1. `@standard_skill_decorators` validates context, starts timing, logs entry
2. `@require_params` checks for 'dataset' parameter
3. `@retry` wraps execution with retry logic
4. `@cached` checks cache before execution
5. `@enrich_metadata` adds metadata to result

## Type Safety Guarantees

MyPy Validation

```
# Type checking passes with decorators
mypy core/decorators.py skills/advanced_analytics.py
# Success: no issues found
```

## IDE Support

```
@timed
def execute(context: AgentContext) -> AgentResult:
    return AgentResult(...)

# IDE knows:
# - Function name: execute
# - Parameter: context: AgentContext
# - Return type: AgentResult
# - Autocomplete works perfectly
```

## Runtime Type Checking

```
from typing import runtime_checkable

# Decorators preserve runtime behavior
assert callable(execute)
assert execute.__name__ == "execute"
assert execute.__annotations__ == {
    'context': AgentContext,
    'return': AgentResult
}
```

## Performance Impact

### Decorator Overhead

Decorator	Overhead	Notes
@validate_context	~0.1ms	Type checks only
@timed	~0.05ms	perf_counter calls
@logged	~0.5ms	Logging I/O
@cached	~0.1ms (hit)	Dict lookup
@cached	~0ms (miss)	Executes normally
@retry	0ms (success)	No overhead on success
@require_params	~0.1ms	Dict key checks
@enrich_metadata	~0.05ms	Dict update

**Total overhead for full stack:** ~1-2ms

**Trade-off:** Minimal overhead for significant functionality

## Best Practices

DO:

- Stack decorators logically** (validation → execution → post-processing)
- Use @standard\_skill\_decorators for consistency**
- Add caching for expensive operations**
- Add retry for flaky external services**
- Validate parameters early with @require\_params**

DON'T:

- Don't cache non-deterministic operations**
- Don't retry operations with side effects**

- ✖ Don't stack same decorator multiple times
- ✖ Don't use decorators for business logic
- ✖ Don't forget to test decorated functions

## Testing Decorated Functions

```
import pytest
from core.decorators import timed, cached
from core.protocols import AgentContext, ResultStatus

def test_decorated_skill():
    """Test that decorators don't break functionality."""

    @timed
    def simple_skill(context: AgentContext) -> AgentResult:
        return AgentResult(
            status=ResultStatus.SUCCESS,
            data={"value": 42},
            message="Success"
        )

    # Test execution
    context = AgentContext(task="test")
    result = simple_skill(context)

    # Verify core functionality
    assert result.success is True
    assert result.data["value"] == 42

    # Verify decorator added metadata
    assert "execution_time_ms" in result.metadata
    assert isinstance(result.metadata["execution_time_ms"], float)
```

## Creating Custom Decorators

```
from typing import Callable, TypeVar, Any, cast
from functools import wraps

F = TypeVar("F", bound=Callable[..., Any])

def my_custom_decorator(func: F) -> F:
    """
    Custom decorator with type safety.

    Template for creating new decorators that preserve types.
    """

    @wraps(func)
    def wrapper(*args: Any, **kwargs: Any) -> Any:
```

```
# Pre-execution logic
print(f"Before {func.__name__}")

# Execute function
result = func(*args, **kwargs)

# Post-execution logic
if isinstance(result, AgentResult):
    result.metadata["custom_decorator"] = True

print(f"After {func.__name__}")
return result

# Cast preserves type information
return cast(F, wrapper)

# Usage
@my_custom_decorator
def execute(context: AgentContext) -> AgentResult:
    return AgentResult(...) # Type safety maintained!
```

## Advanced: Parameterized Decorators

```
def my_parameterized_decorator(
    param1: str,
    param2: int = 10
) -> Callable[[F], F]:
    """
    Decorator factory that accepts parameters.

    Returns a decorator that uses the provided parameters.
    """

    def decorator(func: F) -> F:
        @wraps(func)
        def wrapper(*args: Any, **kwargs: Any) -> Any:
            # Use param1 and param2 here
            print(f"Decorator params: {param1}, {param2}")
            return func(*args, **kwargs)

        return cast(F, wrapper)

    return decorator

# Usage
@my_parameterized_decorator("custom_value", param2=20)
def execute(context: AgentContext) -> AgentResult:
    return AgentResult(...)
```

## Comparison with Other Approaches

Approach	Type Safety	IDE Support	Runtime Cost	Flexibility
TypeVar Decorators	✓✓✓	✓✓✓	Low	High
Traditional Decorators	X	X	Low	High
Inheritance	✓✓	✓✓	None	Low
Middleware Pattern	✓	✓	Medium	High
Aspect-Oriented	✓	✓	High	Medium

## Integration with Framework

Decorators integrate seamlessly with the architecture:

```
# In skills/my_skill.py

from core.decorators import standard_skill_decorators, cached
from core.protocols import AgentContext, AgentResult

@cached(ttl_seconds=300)
@standard_skill_decorators
def execute(context: AgentContext) -> AgentResult:
    """
    Skill with automatic:
    - Context validation
    - Timing
    - Logging
    - Caching
    """
    return AgentResult(...)

# Registry discovers it
# Orchestrator executes it
# Decorators enhance it
# All with perfect type safety!
```

## Summary

TypeVar-based decorators provide:

1. **Type Safety:** Preserves function signatures for static analysis
2. **IDE Support:** Full autocomplete and type checking
3. **Modularity:** Compose concerns independently
4. **Performance:** Minimal overhead (~1-2ms)
5. **Flexibility:** Easy to add, remove, or customize
6. **Testing:** Decorators don't interfere with tests

**Philosophy:** "Cross-cutting concerns should be composable, type-safe, and invisible to domain logic."

The decorator system exemplifies modern Python: leveraging advanced type features (TypeVar, ParamSpec, cast) to build robust, maintainable systems without sacrificing developer experience.