

# D<sup>4</sup> Domain-Driven Database Design

---

## Curriculum Development Discussion Outline

**Presenter:** Peter Heller, Database Systems Instructor, Queens College CUNY

**Purpose:** Exploratory discussion to assess alignment and interest in curriculum development funding

---

### 1. Executive Overview (2-3 minutes)

#### The Problem

- Industry gap: 70%+ of database projects fail or significantly exceed budgets
- Modern developers lack systematic training in database design fundamentals and in-depth analysis skills
- Industry increasingly requires developer/analyst hybrid roles—not separated developers and business analysts
- Disconnect between application domain models (DDD) and physical database implementation
- Students graduate without practical experience in full-lifecycle database design

#### The Solution: D<sup>4</sup> (Domain-Driven Database Design)

- Physical Data Model-first methodology for operational databases
- Integrates Business Glossary Definitions (BGD) into Fully Qualified Domain Names (FQDN)
- Semantic naming conventions enable each table to produce unique golden records, minimizing MDM overhead
- Governance-as-code approach using standardized enterprise object naming
- Practical, teachable framework refined since 2003 (NYC DCAS EC3 energy billing system)

#### The Opportunity

Scaffolded multi-course curriculum where students learn methodology while delivering solutions architect-level project work

---

### 2. Curriculum Scaffold Approach

#### **Core Philosophy: Solutions Architect Training Through Practical Projects**

Students learn D<sup>4</sup> methodology by working through progressively complex database projects—the same deliverables a solutions architect would produce in industry. Each course builds on the previous, creating a coherent multi-semester pathway.

---

#### **Course Sequence Structure**

##### **Course 1: Database Systems Foundations (CSCI 331 equivalent)**

*Prerequisites: Data Structures, Programming proficiency*

### **Learning Objectives:**

- Master Physical Data Model-first design approach
- Implement FQDN taxonomy with BGD integration
- Develop semantic naming conventions for enterprise standardization
- Design schemas where each table produces unique golden records
- Understand governance-as-code principles

### **Practical Project Deliverables:**

- Physical database design for operational system (e.g., order management, student registration)
- ANSI SQL DDL scripts with semantic naming enforcement
- Documentation showing BGD → FQDN mapping
- Normalization analysis demonstrating SRP at table level

### **Technology Stack:**

- SQL Server 2025 (ANSI SQL compliant)
- Docker for consistent dev environments
- GitHub for version control and collaboration
- Azure Data Studio or SQL Server Management Studio

**Student Role:** Database Designer/Developer

---

## **Course 2: Advanced Database Design & Analysis (New or CSCI 381 Enhanced)**

*Prerequisites: Course 1 (Database Systems Foundations)*

### **Learning Objectives:**

- Apply D<sup>4</sup> methodology to complex business domains
- Integrate analysis skills with database design (unified developer/analyst role)
- Work with LLM APIs for schema generation and validation
- Optimize for both transactional and analytical workloads
- Create data models that support enterprise MDM strategies

### **Practical Project Deliverables:**

- Enterprise-scale database design (e.g., financial system, healthcare application)
- Requirements analysis documentation showing business rules → FQDN translation
- AI-assisted schema validation using LLM API calls (via Ollama for cost containment)
- Data lineage documentation demonstrating golden record production
- Migration strategy from legacy schema to D<sup>4</sup> compliant design

### **Technology Stack:**

- SQL Server 2025 with advanced features
- Ollama for local LLM deployment (student cost containment)

- Docker Compose for multi-container environments
- GitHub with CI/CD pipelines
- Erwin Data Modeler or similar for visual documentation

**Student Role:** Database Solutions Architect/Analyst

---

### Course 3: Production Database Engineering (New or CSCI 531 Enhanced)

*Prerequisites: Course 2 (Advanced Database Design & Analysis)*

#### Learning Objectives:

- Deploy D<sup>4</sup> designs to production-grade environments
- Implement monitoring and observability for database systems
- Apply governance-as-code throughout full lifecycle
- Integrate databases with microservices architectures
- Handle real-world data migration and refactoring scenarios

#### Practical Project Deliverables:

- Dockerized database deployment with orchestration
- Automated schema validation and testing pipeline
- Performance optimization demonstrating FQDN benefits
- API integration layer following domain-driven principles
- Production readiness assessment and documentation

#### Technology Stack:

- Docker Swarm or Kubernetes for orchestration
- GitHub Actions for CI/CD automation
- Monitoring tools (Prometheus, Grafana)
- API frameworks integrated with database layer
- Infrastructure-as-code (Terraform or similar)

**Student Role:** Production Database Engineer/DevOps

---

### Scaffolding Principles Across Courses

1. **Progressive Complexity:** Course 1 = single domain, Course 2 = enterprise integration, Course 3 = production deployment
2. **Consistent Methodology:** D<sup>4</sup> principles reinforced in each course with increasing sophistication
3. **Real Deliverables:** Students build portfolio-quality work that demonstrates solutions architect capabilities
4. **Cost-Conscious Technology:** Ollama for AI features eliminates per-student API costs
5. **Industry-Ready Tools:** Docker, GitHub, ANSI SQL ensure transferable skills

---

### Key Differentiator: FQDN Integration

## What Makes This Unique:

- Physical Model-first approach naturally integrates Business Glossary Definitions into table/column names
- FQDN taxonomy creates self-documenting schemas
- Each table designed to produce golden records, reducing enterprise MDM complexity
- Governance enforced at design time through semantic naming, not post-deployment reviews

## Example:

```
Bad: customer_orders, order_items  
Good: Customer.Order, Customer.Order.Item
```

The FQDN immediately shows:

- Domain context (Customer)
- Entity hierarchy (Order contains Items)
- Business glossary mapping (Customer, Order, Item are enterprise terms)
- Golden record capability (Customer.Order is authoritative for order data)

## 3. Core Technical Stack (Standardized Across Courses)

### Rationale: Industry-Standard, Cost-Effective, Transferable Skills

#### Database Platform:

- SQL Server 2025 (ANSI SQL compliant, widely used in enterprise)
- Justification: Students learn transferable SQL skills, strong tooling ecosystem

#### Development Environment:

- Docker for containerization (consistent environments across student machines)
- GitHub for version control, collaboration, and portfolio building
- Azure Data Studio or SSMS (free, professional-grade tools)

#### AI Integration:

- Ollama for local LLM deployment (eliminates per-student API costs)
- Schema validation, documentation generation, requirements analysis support
- Students learn AI-assisted development without financial barriers

#### Infrastructure-as-Code:

- Docker Compose for multi-container orchestration
- GitHub Actions for CI/CD pipelines
- Keeps focus on database design while teaching modern DevOps practices

#### Justification for Choices:

- All tools are free or open-source (no student financial burden)

- SQL Server 2025 available through Azure for Students program
  - Docker/GitHub are industry expectations for any developer role
  - Ollama enables AI features without recurring costs
- 

## 4. Impact Metrics (Focused & Measurable)

### Student Learning Outcomes

1. **Design Competency:** Can students create Physical Data Models that integrate BGD into FQDN?
  - Assessment: Capstone project evaluated by industry advisory board
2. **Developer/Analyst Integration:** Do students demonstrate unified analysis + implementation skills?
  - Assessment: Project documentation showing requirements → schema translation
3. **Governance Understanding:** Can students explain and enforce semantic naming conventions?
  - Assessment: Code reviews, peer evaluations of schema quality

### Employment Outcomes

1. **Job Placement:** Percentage of students securing database-related roles within 6 months
2. **Industry Feedback:** Employer surveys on graduate preparedness
3. **Portfolio Quality:** Student GitHub repositories showcasing D<sup>4</sup> projects

### Curriculum Effectiveness

1. **Course Completion:** Retention rates across three-course sequence
2. **CUNY Adoption:** Number of other CUNY campuses piloting D<sup>4</sup> curriculum
3. **Open Resource Usage:** Downloads/forks of curriculum materials on GitHub

### Long-Term Impact

1. **Technical Debt Reduction:** Track alumni reporting schema quality improvements in their organizations
2. **Methodology Adoption:** Companies implementing D<sup>4</sup> principles after hiring graduates
3. **Community Growth:** Contributions to D<sup>4</sup> methodology from students and alumni

**Key Principle:** Metrics focus on student preparedness for developer/analyst hybrid roles, not research citations or tool downloads

## 5. Peter's Unique Qualifications

### Teaching Excellence

- **10 years teaching** CSCI 331 (Database Systems) and CSCI 381 (Erwin Data Modeling, Power BI) at Queens College
- Developed and refined D<sup>4</sup> methodology starting in **2003** with NYC DCAS EC3 energy billing system
- FQDN taxonomy originated from that 2003 project—20+ years of real-world validation
- Created comprehensive educational materials integrating modern tools (Python, Docker, Pydantic)
- Student-centered pedagogy emphasizing solutions architect-level deliverables

## Methodology Origins & Evolution

- **2003:** Designed NYC DCAS EC3 energy billing system using Fully Qualified Domain Name taxonomy
- **Foundation:** FQDN approach became the basis for D<sup>4</sup> Domain-Driven Database Design
- **Refinement:** 20+ years of iteration across enterprise projects and classroom teaching
- **Validation:** Methodology proven in both production systems and educational settings

## Industry Experience

- Real-world database architecture for municipal government systems
- Enterprise-scale projects requiring MDM integration and golden record production
- Solutions architect perspective: requirements analysis through production deployment
- Bridge between academic rigor and industry pragmatism

## Technical Infrastructure

- Production-grade development environment (multi-node Docker Swarm cluster)
  - Expertise across SQL Server, PostgreSQL, DuckDB
  - Modern development stack supporting AI-integrated workflows
  - Can provide students with cloud-like environment using local infrastructure
- 

## 6. Strategic Partnerships & Industry Connections

### New York CEO Jobs Council

- **Critical Connection:** Peter was in the third cohort
- Direct access to NYC tech employers seeking database-skilled graduates
- Industry advisory board potential from council network
- Employment pipeline for students completing D<sup>4</sup> curriculum
- Real-world validation through employer feedback

### Potential Collaborators

- **CUNY System:** Multi-campus pilot across undergraduate CS programs
- **Microsoft:** SQL Server 2025 education initiatives, Azure for Students integration
- **NYC Tech Employers:** Internship placements, capstone project sponsors
- **Database Vendors:** Erwin, Snowflake, Databricks for tooling partnerships
- **Professional Organizations:** ACM, IEEE database education committees

### Industry Advisory Board (To Be Formed)

- NYC CEO Jobs Council members
- Database architects from financial services, healthcare, government
- Hiring managers seeking developer/analyst hybrid talent
- Alumni working in database-related roles

### Value Proposition to Partners:

- Access to pipeline of students with rare combination: design + analysis + governance skills

- Opportunity to shape curriculum based on real industry needs
  - Early engagement with graduates trained in modern database practices
- 

## 7. Discussion Questions for Funder

### Understanding Their Priorities

1. Does this scaffolded curriculum approach align with your educational funding priorities?
2. What evidence/metrics would be most compelling for your review process?
3. Are there specific gaps in database education you've identified that this addresses?
4. Would multi-institutional collaboration (across CUNY) strengthen the proposal?
5. Is the developer/analyst hybrid role something you see as increasingly important in industry?

### Addressing Potential Concerns

- "**How is this different from existing database courses?**"
  - Physical Model-first with FQDN/BGD integration
  - Developer + analyst skills taught together, not separately
  - Solutions architect-level project deliverables, not theoretical exercises
- "**Why now? What's changed?**"
  - Industry needs developer/analyst hybrids, not separated roles
  - AI integration (via Ollama) makes advanced tooling accessible to undergraduates
  - Modern DevOps practices (Docker, GitHub) now expected for all developers
- "**What about NoSQL/NewSQL?**"
  - D<sup>4</sup> principles (FQDN, semantic naming, golden records) apply across database paradigms
  - Starting with relational provides strongest foundation
- "**How do you ensure adoption beyond Queens College?**"
  - Open curriculum materials via GitHub
  - Instructor guides and workshop for other CUNY faculty
  - NYC CEO Jobs Council connections provide employer demand signal

### Understanding Methodology Status

- D<sup>4</sup> methodology refined over 20+ years (2003 origins)
  - This is curriculum development, not research into untested methodology
  - Focus is on teaching proven practices, not validating new theory
  - Open to discussing intellectual property framework if relevant to funder
- 

## 8. Supporting Materials Available

### Ready to Share:

- D<sup>4</sup> methodology overview document

- Sample curriculum materials from CSCI 331 and CSCI 381
- Student project examples demonstrating D<sup>4</sup> application
- FQDN taxonomy examples from 2003 NYC DCAS EC3 project
- NYC CEO Jobs Council third cohort credentials

#### **Can Be Developed Upon Interest:**

- Detailed course syllabi for three-course sequence
- Assessment rubrics for solutions architect-level deliverables
- Docker/GitHub starter kits for student onboarding
- Ollama integration guides for cost-effective AI features
- Industry advisory board composition and commitment letters

## 9. Contact Information

### **Peter Heller**

Database Systems Instructor, Queens College CUNY

Email: peter.heller@qc.cuny.edu

GitHub: [github.com/QCadjunct](https://github.com/QCadjunct) (educational), [github.com/ph3ll3r](https://github.com/ph3ll3r) (personal)

---

## Summary: Why This Curriculum Matters

### **The Gap:**

- Employers need developer/analysts who can design, implement, and maintain databases
- Current education separates these skills or treats database design as theoretical exercise
- Students graduate without portfolio-quality database work

### **The Solution:**

- Three-course scaffold teaching D<sup>4</sup> methodology through practical projects
- Students produce solutions architect-level deliverables
- FQDN/BGD integration creates self-documenting, governance-compliant schemas
- Cost-effective AI integration via Ollama
- Industry connections via NYC CEO Jobs Council ensure employment outcomes

### **The Track Record:**

- 20+ years of methodology refinement (2003 NYC DCAS EC3 origins)
- 10 years classroom validation at Queens College
- Proven in both production systems and educational settings
- Ready to scale across CUNY system with funder support