

D⁴ Curriculum + AI-Driven Software Development

Integration Outline: Teaching Database Design in the Age of Agentic AI

Context: This outline connects the D⁴ (Domain-Driven Database Design) curriculum scaffold with emerging concepts from "Demystifying AI-Driven Software Development: From Blueprints to Autonomous Builders" by Heller and Maric (Jan 2024).

Executive Summary

The Convergence:

- AI agents (Claude, Anthropic's tools) are transforming software development workflows
- Specification-Driven Development (SDD) + Agentic CLI Programming enable LLMs to generate code from precise specifications
- **D⁴ methodology provides the "blueprint" (physical data model) that AI agents need to generate database implementations**
- Students learn to create specifications that both humans AND AI can execute

Why This Matters for D⁴ Curriculum:

- Database design becomes the **blueprint** in SDD workflow
 - FQDN taxonomy provides the **structured specifications** AI needs
 - Students become "specification architects" who direct AI agents
 - Practical validation: Ollama + LangGraph can implement D⁴ designs automatically
-

Part 1: The Blueprint is the Database Schema

Specification-Driven Development (SDD) Applied to Databases

From the Article:

"SDD is like having that recipe: define the software's behavior, structure, and rules using structured formats"

D⁴ Connection: The Physical Data Model (with FQDN taxonomy) IS the structured specification:

Traditional Approach:

Requirements → Logical Model → Physical Model → Implementation

D⁴ + AI Approach:

Requirements → Physical Model (D⁴/FQDN) → AI Agent Implements → Validation

↑

The Blueprint

What This Means for Students:

1. Course 1 Focus: Students learn to create physical models that serve as precise specifications

- FQDN naming = structured format AI can parse
- BGD integration = semantic context AI needs
- Constraints explicitly defined = rules AI enforces

2. Course 2 Focus: Students validate AI-generated implementations against their blueprints

- LLM generates DDL from Physical Model specification
- Students review for governance compliance
- Iterative refinement: human architect + AI builder

3. Course 3 Focus: Students orchestrate AI agents for production deployment

- CI/CD pipelines where AI validates against D⁴ specifications
- Automated schema migration planning via LLM analysis
- Monitoring that checks runtime behavior against blueprint

Part 2: Agentic CLI Programming for Database Tasks

From Article: "An 'agent' AI that doesn't just follow orders but thinks, acts, and fixes things on its own"

Database Architecture with AI Agents:

```

Student as "Solutions Architect"
  ↓
D4 Blueprint (Physical Model)
  ↓
AI Agent Workflow:
  1. Generate DDL
  2. Analyze dependencies
  3. Create test data
  4. Validate constraints
  5. Document relationships
  6. Suggest optimizations

```

Practical Implementation in Curriculum:

Tool Stack Integration:

- **Ollama** (local LLM): Cost-effective AI agent runtime
- **LangGraph**: Workflow orchestration (from article)
- **Claude API**: Advanced reasoning when needed
- **GitHub Actions**: CI/CD with AI validation

Student Workflow (Course 2):

1. **Student creates Physical Model** using D⁴ methodology:

Table: Customer.Order

Columns:

- Customer_Order_ID (PK)
- Customer_ID (FK)
- Order_Date
- Order_Status

Business Rules:

- Order_Status must be in ('Pending', 'Confirmed', 'Shipped', 'Delivered')
- Order_Date cannot be in future

2. Student writes specification for AI agent:

TASK: Generate SQL Server 2025 DDL for Customer.Order table

REQUIREMENTS:

- Follow FQDN naming convention
- Implement check constraints for Order_Status
- Add temporal table for audit history
- Generate foreign key to Customer.Customer

OUTPUT: ANSI SQL DDL script

3. AI Agent (via Ollama) generates implementation:

- Parses FQDN structure
- Applies SQL Server best practices
- Includes governance checks
- Generates documentation

4. Student validates output:

- Does DDL match D⁴ blueprint?
- Are naming conventions enforced?
- Is temporal tracking implemented correctly?
- Does it produce golden records as designed?

5. Iterative refinement:

- Student provides feedback to AI
- AI regenerates with corrections
- Process mirrors real-world AI-assisted development

Part 3: Pruning AI Workflows - DAGs for Database Design

From Article: "LangGraph lets you build a graph of connected steps (nodes) and paths (edges)"

Application to Database Design:

Traditional Problem:

- Students often create circular dependencies in schemas
- Foreign key ordering is complex
- Migration scripts fail due to dependency violations

D⁴ + LangGraph Solution:

```
# Conceptual DAG for Database Implementation

class DatabaseBuildDAG:
    """
    AI agent analyzes D4 Physical Model and creates dependency graph
    """

    nodes = {
        "analyze_fqdn": "Parse FQDN hierarchy from Physical Model",
        "identify_domains": "Extract domain boundaries (Customer, Product, Order)",
        "order_tables": "Create topological sort of tables by FK dependencies",
        "generate_ddl": "Create DDL respecting dependency order",
        "validate_naming": "Check FQDN compliance",
        "test_golden_records": "Verify each table produces unique records",
        "generate_docs": "Create Mermaid diagrams and documentation"
    }

    edges = [
        ("analyze_fqdn", "identify_domains"),
        ("identify_domains", "order_tables"),
        ("order_tables", "generate_ddl"),
        ("generate_ddl", "validate_naming"),
        ("validate_naming", "test_golden_records"),
        ("test_golden_records", "generate_docs")
    ]
```

Student Learning Outcomes:

- **Course 2:** Students understand how AI decomposes their Physical Model into DAG
- **Course 3:** Students build LangGraph workflows for production deployment
- **Practical Skill:** Students can direct AI agents through complex database tasks

Example Exercise:

"Given this Physical Model with 15 tables spanning 4 domains, create a LangGraph workflow that:

1. Identifies the correct table creation order
2. Generates DDL with proper FK constraints
3. Validates FQDN compliance at each step
4. Stops if any step violates D⁴ principles"

Part 4: CI/CD Pipelines with AI Validation

From Article: "DAGs in CI/CD pipelines... nodes like 'Build code', 'Run tests', 'Deploy'"

D⁴ Database CI/CD with AI Agents:

```
# GitHub Actions Workflow (Conceptual)

name: D4_Schema_Validation

on: [push, pull_request]

jobs:
  validate_d4_compliance:
    runs-on: ubuntu-latest

    steps:
      - name: Checkout schema changes
        uses: actions/checkout@v2

      - name: AI Agent - Parse Physical Model
        uses: ollama-action@v1
        with:
          model: llama3.2
          prompt: |
            Analyze this Physical Model for D4 compliance:
            - Verify FQDN naming conventions
            - Check BGD integration
            - Validate golden record capability

      - name: AI Agent - Generate Test Cases
        run: |
          # LLM generates test data based on constraints
          # Validates each table produces unique golden records

      - name: AI Agent - Detect Schema Drift
        run: |
          # Compare proposed schema to production
          # Identify breaking changes
          # Suggest migration strategy

      - name: Human Review Gate
        if: contains(github.event.head_commit.message, '[BREAKING]')
        run: |
          echo "Breaking change detected - requires manual review"
          # Pause pipeline for architect approval
```

Student Experience (Course 3):

Students learn to:

1. Define what "D⁴ compliant" means in executable form (CI checks)
2. Use AI agents to validate their designs automatically
3. Understand when AI catches issues vs. when human judgment needed
4. Build pipelines that enforce governance-as-code

Key Insight from Article:

"Loose coupling keeps things fast... DAG handles dependencies logically without getting stuck"

Applied to D⁴: Each domain (Customer, Product, Order) can evolve independently because FQDN enforces loose coupling. AI agents can validate changes within one domain without analyzing entire schema.

Part 5: The "Specification Architect" Role

Redefining the Developer/Analyst Hybrid

Traditional Roles:

- **Developer:** Writes code
- **DBA:** Implements database
- **Analyst:** Documents requirements
- **Architect:** Creates designs

New Role (What D⁴ Curriculum Trains):

- **Specification Architect:** Creates precise blueprints that both humans and AI can execute

Skills Students Develop:

Traditional Skill	AI-Enhanced Skill
Write DDL manually	Specify requirements for AI to generate DDL
Debug schema issues	Validate AI-generated schemas against D ⁴ principles
Document designs	Create machine-readable specifications (FQDN)
Manual testing	Define test criteria for AI to implement
Code reviews	Review AI-generated code for governance compliance

From Article:

"Developers focus on the 'what' (the desired result)... AI agents do the 'how' (the code)"

D⁴ Application:

- **Student specifies:** "Customer.Order table following FQDN taxonomy with temporal audit"
 - **AI implements:** DDL, constraints, indexes, test data, documentation
 - **Student validates:** Against D⁴ blueprint and business requirements
-

Part 6: Practical Course Integration

Course 1: Database Systems Foundations + AI Basics

Week 1-4: Traditional D⁴ Concepts

- Physical Model-first design
- FQDN taxonomy
- Semantic naming
- Golden record principles

Week 5-8: Introduction to AI-Assisted Design

- Students install Ollama locally (no cost)
- Basic prompts for DDL generation
- Validating AI output against manual designs
- Understanding when AI helps vs. when it confuses

Week 9-12: Structured Specifications

- Creating "AI-readable" Physical Models
- JSON/YAML representations of D⁴ designs
- Version control for schemas (GitHub)
- Documentation as code

Capstone Project:

Design a database for a small business. Create the Physical Model in D⁴ format, then use Ollama to generate implementation. Compare AI output to your manual implementation. Document where AI succeeded and where it failed.

Course 2: Advanced Design + Agentic Workflows

Week 1-4: Complex Domain Modeling

- Multi-domain schemas (Customer, Product, Order, Inventory)
- FQDN hierarchies across domains
- Cross-domain constraints and golden records
- MDM integration patterns

Week 5-8: LangGraph for Database Workflows

- Installing and configuring LangGraph
- Building DAGs for schema deployment
- Dependency analysis and resolution
- Automated migration planning

Week 9-12: AI-Driven Requirements Analysis

- Students interview "client" (instructor or industry partner)
- Use LLM to extract BGD from unstructured requirements
- Map business terms to FQDN
- Generate Physical Model with AI assistance

Capstone Project:

Enterprise-scale database design (30+ tables across 5 domains). Build LangGraph workflow that:

1. Generates complete DDL from your Physical Model
2. Creates dependency-aware deployment script
3. Generates test data following D⁴ constraints
4. Produces documentation including Mermaid diagrams

Course 3: Production Engineering + AI Operations

Week 1-4: CI/CD for Databases

- GitHub Actions for schema validation
- Automated D⁴ compliance checking
- AI-powered code review for DDL changes
- Blue-green deployments for schema changes

Week 5-8: Observability with AI

- Schema drift detection using LLMs
- Performance analysis and optimization suggestions
- Automated index recommendation
- Query pattern analysis for schema refinement

Week 9-12: AI Governance

- Auditing AI-generated schemas
- Ensuring D⁴ compliance in AI outputs
- Human-in-the-loop workflows
- Ethics of AI-assisted database design

Capstone Project:

Take a legacy schema (provided by instructor). Use AI agents to:

1. Analyze existing schema and map to D⁴ principles
2. Generate migration plan to D⁴-compliant design
3. Build CI/CD pipeline for migration execution
4. Deploy to production-like environment with monitoring
5. Present to industry advisory board

Part 7: Technology Stack Integration

Core Stack (All Courses)

Student Development Environment

Database: SQL Server 2025
Containers: Docker + Docker Compose
Version Control: GitHub
AI Runtime: Ollama (local LLM)
IDE: Azure Data Studio / VS Code

Progressive AI Integration

Course 1 (Basic AI Assistance):

- Ollama with Llama 3.2 (7B model)
- Simple prompts for DDL generation
- Documentation generation
- Basic validation checks

Course 2 (Workflow Orchestration):

- LangGraph for multi-step tasks
- Claude API for complex reasoning (via Anthropic for Students)
- Requirements analysis and BGD extraction
- Automated diagram generation

Course 3 (Production AI Operations):

- GitHub Actions with AI agents
- Monitoring and alerting
- Schema evolution management
- Automated optimization

Cost Control:

- Ollama runs locally (free, no API costs)
- Anthropic provides educational credits for Claude API
- SQL Server 2025 via Azure for Students (free)
- All other tools open-source or free tier

Part 8: Measuring Success in AI Era

Updated Learning Outcomes

Course 1 Assessment:

- Can student create D⁴-compliant Physical Model?
- Can student validate AI-generated DDL against their design?
- Does student understand when to trust AI vs. manual implementation?

Course 2 Assessment:

- Can student build LangGraph workflow for complex database task?

- Can student extract BGD from unstructured requirements using LLM?
- Does student produce enterprise-quality specifications for AI execution?

Course 3 Assessment:

- Can student deploy CI/CD pipeline with AI validation?
- Can student audit AI-generated schemas for governance compliance?
- Does student demonstrate human-AI collaboration best practices?

Employment Outcomes (AI-Enhanced Roles)

Students prepared for:

- **AI-Native Database Developer:** Uses AI for implementation, focuses on specification quality
- **MLOps Engineer (Database Focus):** Builds AI workflows for database operations
- **Data Platform Architect:** Designs blueprints for AI-assisted implementation
- **DevOps Engineer (Database Specialization):** Automates database lifecycle with AI

NYC CEO Jobs Council Validation: These are the roles employers are seeking—developers who can work WITH AI, not be replaced BY AI.

Part 9: Research Questions This Enables

Empirical Studies We Can Conduct

1. AI vs. Human Implementation Quality:

- Same Physical Model → Human implementation vs. AI implementation
- Measure: correctness, completeness, performance, maintainability

2. Specification Quality Impact:

- Well-defined FQDN vs. ambiguous naming → AI output quality
- Measure: how often AI produces D⁴-compliant schemas

3. Learning Curve Comparison:

- Traditional teaching vs. AI-assisted learning
- Measure: time to competency, error rates, confidence levels

4. Governance Enforcement:

- Manual code review vs. AI-powered validation
- Measure: compliance rates, false positive/negative rates

5. Developer Productivity:

- Solo work vs. human-AI collaboration
- Measure: tasks completed, quality metrics, job satisfaction

Publishable Findings:

- "Teaching Database Design in the Age of AI Agents"
 - "FQDN Taxonomy as Machine-Readable Specifications for LLMs"
 - "Human-AI Collaboration in Database Architecture"
 - "From Code Writers to Specification Architects: Educational Transformation"
-

Part 10: Competitive Advantages

Why D⁴ + AI Curriculum is Unique

What Others Are Doing:

- Traditional database courses: Manual DDL writing, no AI integration
- "AI for databases" courses: Focus on using AI to query data, not design schemas
- Data science programs: Skip database design entirely, use NoSQL/DataFrames

What We're Doing Differently:

1. **Physical Model as AI-Executable Blueprint:** FQDN taxonomy enables AI understanding
2. **Governance-as-Code from Day One:** Students learn to create specifications AI can validate
3. **Solutions Architect Training:** Students direct AI agents, don't just prompt them
4. **40 Years + Cutting Edge:** Proven methodology enhanced with latest AI capabilities
5. **Cost-Effective:** Ollama enables advanced AI features without per-student costs

Market Positioning:

"The only undergraduate database curriculum that teaches students to be specification architects directing AI agents to implement enterprise-quality database systems."

Part 11: Implementation Roadmap

Phase 1: Course 1 Pilot (Semester 1)

Objective: Validate basic D⁴ + AI integration

- Update CSCI 331 syllabus with AI components
- Provide Ollama installation guide for students
- Create 3-5 lab exercises: manual design → AI implementation → validation
- Measure: student engagement, AI output quality, learning outcomes

Deliverables:

- Updated course materials
- Student project portfolio examples
- Assessment data
- Lessons learned report

Phase 2: Full Sequence Development (Year 1)

Objective: Build out Courses 2 and 3

- Design advanced D⁴ + LangGraph curriculum
- Partner with NYC CEO Jobs Council for capstone projects
- Develop industry advisory board for Course 3 evaluations
- Create instructor guides for AI-assisted teaching

Deliverables:

- Three-course sequence fully documented
 - LangGraph workflow templates for students
 - Industry partnership agreements
 - CI/CD pipeline examples
-

Phase 3: CUNY-Wide Pilot (Year 2)

Objective: Scale across CUNY system

- Train 3-5 faculty at other CUNY campuses
- Provide standardized infrastructure (Docker images, GitHub templates)
- Coordinate capstone projects across campuses
- Share results with industry partners

Deliverables:

- Faculty training workshops
 - Multi-campus coordination framework
 - Standardized assessment instruments
 - Adoption metrics and feedback
-

Phase 4: Open Educational Resources (Year 3)

Objective: Maximize impact beyond CUNY

- Publish curriculum materials on GitHub (open-source)
- Create video tutorials and documentation
- Present at conferences (ACM SIGCSE, SIGMOD Education Track)
- Engage with other universities for adoption

Deliverables:

- Complete OER package
 - Conference papers and presentations
 - Community engagement metrics
 - Sustainability plan
-

Part 12: Integration with Heller-Maric Article Concepts

Direct Mappings from Article to D⁴ Curriculum

SDD (Specification-Driven Development)

- **Article Concept:** "Define behavior, structure, rules using structured formats"
- **D⁴ Application:** Physical Data Model with FQDN taxonomy IS the structured specification
- **Student Skill:** Write database blueprints that AI can implement

Agentic CLI Programming

- **Article Concept:** "AI agent that thinks, acts, fixes things on its own"
- **D⁴ Application:** LangGraph workflows for schema deployment
- **Student Skill:** Orchestrate AI agents through complex database tasks

Pruning/Efficiency

- **Article Concept:** "Make AI systems smarter by removing redundant steps"
- **D⁴ Application:** DAG analysis of Physical Model dependencies
- **Student Skill:** Design schemas with minimal coupling, maximum AI efficiency

CI/CD Integration

- **Article Concept:** "Automated testing, building, deploying"
- **D⁴ Application:** GitHub Actions validating D⁴ compliance
- **Student Skill:** Build pipelines enforcing governance-as-code

The Future Vision

- **Article Quote:** "Precise plans, agentic action, pruning, and CI/CD"
- **D⁴ Translation:** Physical Models, AI agents, dependency management, automated deployment
- **Student Outcome:** Graduates who embody this future vision

Summary: The Complete Integration

What This Changes:

1. **D⁴ methodology** (proven since 2003) provides the **structured blueprints** AI needs
2. **Students** become **specification architects** who direct AI implementation
3. **FQDN taxonomy** creates **machine-readable** database designs
4. **Three-course scaffold** progressively builds **human-AI collaboration** skills
5. **Cost-effective AI integration** via Ollama makes advanced features accessible
6. **Industry validation** through NYC CEO Jobs Council ensures relevance

The Pitch to Funders:

"We're not just teaching database design. We're training the next generation of specification architects who can direct AI agents to implement enterprise-quality database systems. Our D⁴ methodology, refined over 20 years, provides the structured blueprints that AI needs to succeed. Students graduate with a rare combination: deep database expertise AND the ability to leverage AI effectively—exactly what NYC employers need."

The Evidence Base:

- **40 years industry experience** (methodology proven in production)
- **10 years classroom validation** (pedagogy tested with real students)
- **2024 cutting-edge AI integration** (article demonstrates current best practices)
- **NYC CEO Jobs Council connection** (direct pipeline to employers)

The Deliverable:

A curriculum that produces graduates who are:

- **Specification architects** (not just code writers)
 - **AI collaborators** (not AI replacements)
 - **Database experts** (with 40 years of methodology)
 - **Job-ready** (validated by NYC employers)
-

Contact for Further Discussion

Peter Heller

Co-author: "Demystifying AI-Driven Software Development" (with Marco Maric)

Database Systems Instructor, Queens College CUNY

Email: peter.heller@qc.cuny.edu

GitHub: github.com/QCAdjunct (educational), github.com/ph3ll3r (personal)

Co-Author:

Marco Maric (Agentic AI and LangGraph expertise)

Appendix: Key Terms Glossary

From D⁴ Methodology:

- **FQDN (Fully Qualified Domain Name):** Hierarchical naming convention for database objects
- **BGD (Business Glossary Definition):** Enterprise-standard terms mapped to FQDN
- **Golden Record:** Authoritative data produced by properly designed table
- **Governance-as-Code:** Automated enforcement of design standards

From AI-Driven Development:

- **SDD (Specification-Driven Development):** Define structure in structured format, AI implements
- **Agentic CLI:** AI agents that act autonomously based on specifications
- **LangGraph:** Framework for building multi-step AI workflows as DAGs
- **Pruning:** Optimizing AI workflows by removing redundant steps

Integration Terms:

- **Specification Architect:** Professional who creates blueprints for AI execution
- **AI-Readable Schema:** Database design with semantic naming AI can parse
- **Human-AI Collaboration:** Working model where humans specify, AI implements, humans validate
- **Governance Validation:** Automated checking of AI output against enterprise standards

